

boston-dataset-example

July 10, 2020

1 Boston Housing Dataset Example

Most of this code is included in Chapter 8 of [Data Science for Mathematicians](#). We convert it to notebook form here so that you can see the output and explore it interactively online yourself.

1.1 Step 1: Obtain data

The Boston housing dataset is built into scikit-learn, so we can import it easily, as follows.

```
[1]: from sklearn.datasets import load_boston
     boston = load_boston()
```

But the `boston` object created this way is a conglomeration of several sub-objects and not ready to be printed in a human-readable way, so we organize it as follows.

1.2 Step 2: Create a feature-target dataset

We extract the features and target into separate variables and inspect their contents.

```
[2]: import pandas as pd
     features = pd.DataFrame(
         data=boston.data,
         columns=boston.feature_names)
     target = boston.target
```

The features are a pandas DataFrame, the first few rows of which are shown here.

```
[3]: features.head()
```

```
[3]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14

2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

How many rows are there, actually?

```
[4]: len( features )
```

```
[4]: 506
```

The target is a NumPy array that can be viewed as another column in the same dataset, as shown here.

```
[5]: target
```

```
[5]: array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9, 15. ,
        18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
        15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
        13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
        21.2, 19.3, 20. , 16.6, 14.4, 19.4, 19.7, 20.5, 25. , 23.4, 18.9,
        35.4, 24.7, 31.6, 23.3, 19.6, 18.7, 16. , 22.2, 25. , 33. , 23.5,
        19.4, 22. , 17.4, 20.9, 24.2, 21.7, 22.8, 23.4, 24.1, 21.4, 20. ,
        20.8, 21.2, 20.3, 28. , 23.9, 24.8, 22.9, 23.9, 26.6, 22.5, 22.2,
        23.6, 28.7, 22.6, 22. , 22.9, 25. , 20.6, 28.4, 21.4, 38.7, 43.8,
        33.2, 27.5, 26.5, 18.6, 19.3, 20.1, 19.5, 19.5, 20.4, 19.8, 19.4,
        21.7, 22.8, 18.8, 18.7, 18.5, 18.3, 21.2, 19.2, 20.4, 19.3, 22. ,
        20.3, 20.5, 17.3, 18.8, 21.4, 15.7, 16.2, 18. , 14.3, 19.2, 19.6,
        23. , 18.4, 15.6, 18.1, 17.4, 17.1, 13.3, 17.8, 14. , 14.4, 13.4,
        15.6, 11.8, 13.8, 15.6, 14.6, 17.8, 15.4, 21.5, 19.6, 15.3, 19.4,
        17. , 15.6, 13.1, 41.3, 24.3, 23.3, 27. , 50. , 50. , 50. , 22.7,
        25. , 50. , 23.8, 23.8, 22.3, 17.4, 19.1, 23.1, 23.6, 22.6, 29.4,
        23.2, 24.6, 29.9, 37.2, 39.8, 36.2, 37.9, 32.5, 26.4, 29.6, 50. ,
        32. , 29.8, 34.9, 37. , 30.5, 36.4, 31.1, 29.1, 50. , 33.3, 30.3,
        34.6, 34.9, 32.9, 24.1, 42.3, 48.5, 50. , 22.6, 24.4, 22.5, 24.4,
        20. , 21.7, 19.3, 22.4, 28.1, 23.7, 25. , 23.3, 28.7, 21.5, 23. ,
        26.7, 21.7, 27.5, 30.1, 44.8, 50. , 37.6, 31.6, 46.7, 31.5, 24.3,
        31.7, 41.7, 48.3, 29. , 24. , 25.1, 31.5, 23.7, 23.3, 22. , 20.1,
        22.2, 23.7, 17.6, 18.5, 24.3, 20.5, 24.5, 26.2, 24.4, 24.8, 29.6,
        42.8, 21.9, 20.9, 44. , 50. , 36. , 30.1, 33.8, 43.1, 48.8, 31. ,
        36.5, 22.8, 30.7, 50. , 43.5, 20.7, 21.1, 25.2, 24.4, 35.2, 32.4,
        32. , 33.2, 33.1, 29.1, 35.1, 45.4, 35.4, 46. , 50. , 32.2, 22. ,
        20.1, 23.2, 22.3, 24.8, 28.5, 37.3, 27.9, 23.9, 21.7, 28.6, 27.1,
        20.3, 22.5, 29. , 24.8, 22. , 26.4, 33.1, 36.1, 28.4, 33.4, 28.2,
        22.8, 20.3, 16.1, 22.1, 19.4, 21.6, 23.8, 16.2, 17.8, 19.8, 23.1,
        21. , 23.8, 23.1, 20.4, 18.5, 25. , 24.6, 23. , 22.2, 19.3, 22.6,
        19.8, 17.1, 19.4, 22.2, 20.7, 21.1, 19.5, 18.5, 20.6, 19. , 18.7,
        32.7, 16.5, 23.9, 31.2, 17.5, 17.2, 23.1, 24.5, 26.6, 22.9, 24.1,
        18.6, 30.1, 18.2, 20.6, 17.8, 21.7, 22.7, 22.6, 25. , 19.9, 20.8,
```

```

16.8, 21.9, 27.5, 21.9, 23.1, 50. , 50. , 50. , 50. , 50. , 13.8,
13.8, 15. , 13.9, 13.3, 13.1, 10.2, 10.4, 10.9, 11.3, 12.3, 8.8,
7.2, 10.5, 7.4, 10.2, 11.5, 15.1, 23.2, 9.7, 13.8, 12.7, 13.1,
12.5, 8.5, 5. , 6.3, 5.6, 7.2, 12.1, 8.3, 8.5, 5. , 11.9,
27.9, 17.2, 27.5, 15. , 17.2, 17.9, 16.3, 7. , 7.2, 7.5, 10.4,
8.8, 8.4, 16.7, 14.2, 20.8, 13.4, 11.7, 8.3, 10.2, 10.9, 11. ,
9.5, 14.5, 14.1, 16.1, 14.3, 11.7, 13.4, 9.6, 8.7, 8.4, 12.8,
10.5, 17.1, 18.4, 15.4, 10.8, 11.8, 14.9, 12.6, 14.1, 13. , 13.4,
15.2, 16.1, 17.8, 14.9, 14.1, 12.7, 13.5, 14.9, 20. , 16.4, 17.7,
19.5, 20.2, 21.4, 19.9, 19. , 19.1, 19.1, 20.1, 19.9, 19.6, 23.2,
29.8, 13.8, 13.3, 16.7, 12. , 14.6, 21.4, 23. , 23.7, 25. , 21.8,
20.6, 21.2, 19.1, 20.6, 15.2, 7. , 8.1, 13.6, 20.1, 21.8, 24.5,
23.1, 19.7, 18.3, 21.2, 17.5, 16.8, 22.4, 20.6, 23.9, 22. , 11.9])

```

1.3 Step 3: Split into training and test datasets

We will use 80% of the data for training and then test our model on the 20% held out for that purpose. Scikit-learn contains a function that will randomly split the dataset for us into training and test sets. We add the `random_state` parameter to specify a random number seed, thus guaranteeing reproducibility of the same results if you re-run this notebook later.

```

[6]: from sklearn.model_selection import train_test_split
      (X_training, X_test, y_training, y_test) = \
          train_test_split(features, target, train_size=0.8, random_state=1)

```

Let's verify that the split produced objects of the appropriate sizes.

```

[7]: len( X_training ), len( X_test ), len( y_training ), len( y_test )

```

```

[7]: (404, 102, 404, 102)

```

Since 404 is almost exactly 80% of the original 506, it looks like this has worked correctly.

1.4 Step 4: Create a model from the training dataset

We use scikit-learn's Pipeline object to compose two steps in sequence: First, select the five best features to use for prediction, and second, use those five features to fit a linear model to the training data.

```

[8]: ### Step 4: Create a model from the training dataset
      from sklearn.pipeline import Pipeline
      from sklearn.feature_selection import SelectKBest
      from sklearn.linear_model import LinearRegression
      estimator = Pipeline([
          ('select', SelectKBest(k=5)),
          ('model', LinearRegression())
      ])
      fit_model = estimator.fit(X=X_training , y=y_training)

```

Let's say we'd like to see which features were selected. We can ask the first step in the pipeline to show us its results.

```
[9]: features.columns[ fit_model[0].get_support() ]
```

```
[9]: Index(['CRIM', 'NOX', 'RM', 'AGE', 'LSTAT'], dtype='object')
```

And if we want to see the coefficients the model assigned to each of those variables, we can ask the second step in the pipeline for its results.

```
[10]: fit_model[1].intercept_, fit_model[1].coef_
```

```
[10]: (2.9524367266279405,  
      array([-0.09549911, -4.08891308,  4.56355544,  0.02161194, -0.61759647]))
```

The resulting model is therefore approximately the following.

$$2.952 - 0.095\text{CRIM} - 4.089\text{NOX} + 4.564\text{RM} + 0.022\text{AGE} - 0.618\text{LSTAT}$$

1.5 Step 5: Score the model using the test set

We compute the root mean squared error of the model on the test set.

```
[11]: predictions = fit_model.predict(X=X_test)  
      from sklearn.metrics import mean_squared_error  
      mean_squared_error(y_test, predictions)**0.5
```

```
[11]: 5.630885425217404
```